Rethink about the Word-level Quality Estimation for Machine Translation from Human Judgement

Zhen Yang, Fandong Meng, Yuanmeng Yan, and Jie Zhou

Pattern Recognition Center, WeChat AI, Tencent Inc, Beijing, China {zieenyang, fandongmeng, withtomzhou}@tencent.com

Abstract

Word-level Quality Estimation (QE) of Machine Translation (MT) aims to find out potential translation errors in the translated sentence without reference. Typically, conventional works on word-level QE are designed to predict the translation quality in terms of the post-editing effort, where the word labels ("OK" and "BAD") are automatically generated by comparing words between MT sentences and the post-edited sentences through a Translation Error Rate (TER) toolkit. While the post-editing effort can be used to measure the translation quality to some extent, we find it usually conflicts with the human judgement on whether the word is well or poorly translated. To overcome the limitation, we first create a golden benchmark dataset, namely HJQE (Human Judgement on Quality Estimation), where the expert translators directly annotate the poorly translated words on their judgements. Additionally, to further make use of the parallel corpus, we propose the self-supervised pre-training with two tag correcting strategies, namely tag refinement strategy and tree-based annotation strategy, to make the TER-based artificial QE corpus closer to HJQE. We conduct substantial experiments based on the publicly available WMT En-De and En-Zh corpora. The results not only show our proposed dataset is more consistent with human judgment but also confirm the effectiveness of the proposed tag correcting strategies.¹

1 Introduction

Quality Estimation (QE) of Machine Translation (MT) aims to automatically estimate the quality of the translation generated by MT systems, with no reference available. It typically acts as a post-processing module in commercial MT systems, alerting the user with potential translation errors. Figure 1 shows an example of QE, where the

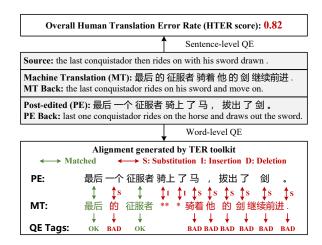


Figure 1: The illustration of the sentence-level and word-level QE tasks. The word-level QE tags are generated by the TER toolkit.

sentence-level task predicts a score indicating the overall translation quality, and the word-level QE needs to predict each word as OK or BAD^2 .

Currently, the collection of QE datasets mainly relies on the Translation Error Rate (TER) toolkit (Snover et al., 2006). Specifically, given the machine translations and their corresponding postedits (PE, generated by human translators) or target sentences of parallel corpus as the pseudo-PE (Tuan et al., 2021; Lee, 2020), the rule-based TER toolkit is used to generate the word-level alignment between the MT and the PE based on the principle of minimal editing. All MT words not aligned to PE are annotated as BAD (shown in Figure 1). Such annotation is also referred as post-editing effort (Fomicheva et al., 2020; Specia et al., 2020). The post-editing effort measures the translation quality in terms of the efforts the translator need to spend to transform the MT sentence to the golden reference. However, we find it usually conflicts with human judgements on whether the word is well or

¹The data can be found at https://github.com/ ZhenYangIACAS/HJQE.

²In this paper, we mainly focus on the word-level QE on the target side, which aims to detect potential translation errors in MT sentences.

Source: It is happy for me to be asked to speak here.

MT: <u>我 很 高兴</u> 被 <u>要求</u> 在 这里 <u>发言</u> 。 MT Back: <u>I am so happy</u> to be <u>asked</u> to <u>speak</u> here.

PE: 被 邀请 在 这里 讲话 我 很 高兴 。 PE Back: Being invited to talk here makes me so happy.

TER-based: **我 很 高兴** 被 **要求** 在 这里 <mark>发言</mark> 。 Human: 我 很 高兴 被 要求 在 这里 发言 。

a) Some words in MT are mistakenly annotated to BAD though the overall semantic is not changed.

Source: The Zaporizhian Hetman was then dispatched to Istanbul, and impaled on hooks.

MT: 扎 波罗 齐安海 特曼 号 随后 被 派 往 伊斯坦布尔, 并 被 撞 在 钩 上。

MT Back: The Zaporizhian Hetman was then dispatched to Istanbul, and was bumped on the hook

PE: Zaporizhian Hetman 随后 被 派 往 伊斯坦布尔 , 并 被 钉 在 钩子 上 。

PE Back: Zaporizhian Hetman was then dispatched to Istanbul, and was nailed on hooks.

TER-based: 扎波罗齐安海特曼号随后被派往伊斯坦布尔,并被撞在钩上。

Human: 扎 波罗 齐安海 特曼 号 随后 被 派 往 伊斯坦布尔 , 并 被 撞 在 钩 上 。

b) Human annotate the clause "被撞在钩上" as a whole, while TER-based annotations are fragmented.

Figure 2: Two examples show the gap between the TER-based annotation and human's direct annotation on detecting translation errors. The red color indicates BAD tags (text with translation errors), while the green color indicates OK tags (text without translation errors).

poorly translated. There are two main issues that result in the conflicts. First, the PE sentences often substitute some words with better synonyms and reorder some sentence constituents for polish purposes. These operations do not destroy the meaning of the translated sentence, but make some words mistakenly annotated under the exact matching criterion of TER (shown in Figure 2a). Second, when fatal errors occur in MTs, a human annotator typically takes the whole sentence or clause as BAD. However, the TER-based annotation still try to find trivial words that align with PE, resulting in fragmented annotations (shown in Figure 2b). In many application scenarios and down-stream tasks, it is usually important even necessary to detect whether the word is well or poorly translated from the human judgement (Yang et al., 2021). However, most previous works still use the TER-based dataset as the evaluation benchmark, which makes the models' predictions deviate from the human judgement.

To overcome the limitations stated above, for the first time, we concentrate on the model's ability of finding translation errors on MT sentences from the human judgement. We first collect a high quality benchmark dataset *HJQE* where human annotators directly annotate the text spans that lead to the translation errors in MT sentences. Then, to further make use of the large scale translation parallel corpus, we also propose two tag correcting strategies, namely tag refinement strategy and tree-based annotation strategy, which make the TER-based annotations more consistent with human judgment.

Our contributions can be summarized as follows: 1) We collect a new dataset called *HJQE*



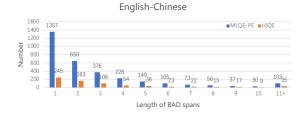


Figure 3: The length distribution of BAD spans on English-German and English-Chinese validation sets.

that directly annotates translation errors on MT sentences. We conduct detailed analyses and demonstrate two differences between *HJQE* and the previous TER-based dataset. 2) To make use of the large scale translation parallel corpus, we propose self-supervised pre-training approach with two automatic tag correcting strategies to make the TER-based artificial dataset more consistent with human judgment and then boost the performance by large-scale pre-training. 3) We conduct experiments on our collected *HJQE* dataset as well as the TER-based dataset MLQE-PE. Experimental results of the automatic and human evaluation show that our approach achieves higher consistency with human judgment.

Dataset	Split		English-German				English-Chinese				
		samples	tokens	MT BAD tags	MT Gap BAD tags	samples	tokens	MT BAD tags	MT Gap BAD tags		
MLQE-PE	train	7000	112342	31621 (28.15%)	5483 (4.59%)	7000	120015	65204 (54.33%)	10206 (8.04%)		
	valid	1000	16160	4445 (27.51%)	716 (4.17%)	1000	17063	9022 (52.87%)	1157 (6.41%)		
HJQE (ours)	train	7000	112342	10804 (9.62%)	640 (0.54%)	7000	120015	19952 (16.62%)	348 (0.27%)		
	valid	1000	16160	1375 (8.51%)	30 (0.17%)	1000	17063	2459 (14.41%)	8 (0.04%)		
	test	1000	16154	993 (6.15%)	28 (0.16%)	1000	17230	2784 (16.16%)	11 (0.06%)		

Table 1: Statistics of TER-based MLQE-PE dataset and our proposed HJQE dataset.

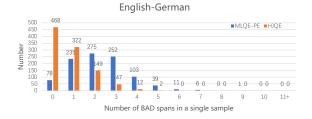
2 Data Collection and Analysis

2.1 Data Collection

To make our collected dataset comparable to TER-generated ones, we directly take the source and MT texts from MLQE-PE (Fomicheva et al., 2020), the official dataset for the WMT20 QE shared task. It includes two language pairs that contain TER-generated annotations: English-German (En-De) and English-Chinese (En-Zh). The source texts are sampled from Wikipedia documents and the translations are obtained from the Transformer-based neural machine translation (NMT) system (Vaswani et al., 2017).

Our data collection follows the following process. First, we hire a number of translator experts. where 5 translators for En-Zh and 6 for En-De. They are all graduated students that major in the translation and have the professional ability on the corresponding translation direction. For each sample, we randomly distribute it to two annotators. Each annotator is provided only the source sentence and its corresponding translation (but without the context or passage which the source sentence is taken from). For En-Zh, the translations are tokenized (as they are in MLQE-PE). Note that although the PE sentences exist in MLQE-PE, the human annotators have no access to them, making the annotation process as fair and unbiased as possible. After one sample is both annotated by the two annotators, we check whether the annotations are consistent. If they have annotation conflicts, we will re-assign the sample to other two annotators until we get the consistent annotations.

For the annotation protocol, we ask human translators to find words, phrases, clauses or even the whole sentences that contain translation error in MT sentences, and annotate them as BAD tags. Here, the translation error means the translation distorts the meaning of the source sentence, but excluding minor mismatches such as synonyms and punctuation. Meanwhile, if the translation does not conform to the grammar of the target language,



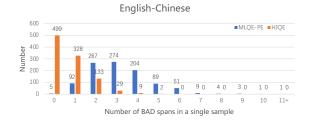
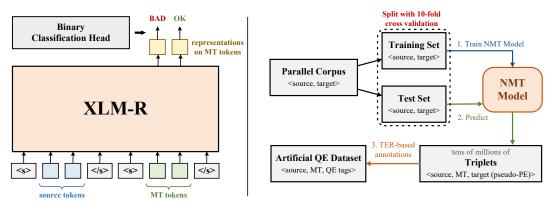


Figure 4: The distribution that reveals how many BAD spans in every single validation sample.

they should also find them and annotate as BAD. The annotation and distribution of samples are automatically conducted through the annotation system. After all samples are annotated, we ask another translator (1 for En-Zh and 1 for En-De, and they do not participant in the annotation process), sampling a small proportion (400 samples) of the full annotated dataset and ensure the accuracy is above 98%.

2.2 Statistics and Analysis

Overall Statistics. In Table 1, we show detailed statistics of the collected *HJQE*. For comparison, we also present the statistics of MLQE-PE which is automatically annotated with TER. First, we see that the total number of BAD tags decreases heavily when human's annotations replaces the TER-based annotations (from 28.15% to 9.62% for En-De, and from 54.33% to 16.62% for En-Zh). It indicates that the human's annotations tends to annotate OK as long as the translation correctly expresses the meaning of the source sentence, but ignores the secondary issues like synonym substitutions and constituent reordering. Second, we find the number of BAD tags in the gap (indicating a few words



- a) The overall architecture of our model.
- b) The construction steps of artificial QE dataset for pre-training.

Figure 5: The model architecture and the construction of artificial QE dataset.

are missing between two MT tokens) also greatly decreases. It's because that human's annotations tends to regard the missing translations (i.e., the BAD gaps) and the translation errors as a whole but only annotate BAD tags on MT tokens 3 .

The Length of BAD Spans. We show the number of BAD spans⁴ of different lengths in Figure 3. We can see that most BAD spans only contain a few tokens, showing the well-known long-tail distribution. For En-De, the long-tail distribution is sharper, where 70.5% of BAD spans are one-token spans. When comparing the TER-based annotations with human's, we find that human's annotation includes fewer BAD spans of each length, but the overall distribution is similar.

Unity of BAD Spans. To reveal the unity of the human's annotations, we group the samples according to the number of BAD spans in each single sample, and show the overall distribution. From Figure 4, we can find that the TER-based annotations follow the Gaussian distribution, where a large proportion of samples contain 2, 3, or even more BAD spans, indicating the TER-based annotations are fragmented. However, our collected annotations on translation errors are more unified, with only a small proportion of samples including more than 2 BAD spans. Besides, we find a large number of samples that are fully annotated as OK in human's annotations. However, the number is extremely small for TER-based annotations (78 in English-German and 5 for English-Chinese). This shows a large proportion of BAD spans in TERbased annotations do not really destroy the semantic of translations and are thus regarded as OK by human annotators.

3 Approach

In this section, we will first introduce the backbone of the model and the self-supervised pre-training approach based on the large scale parallel corpus. Then, we propose two correcting strategies to make the TER-based artificial tags closer to the human judgment.

3.1 Model Architecture

Following (Ranasinghe et al., 2020; Lee, 2020; Moura et al., 2020; Ranasinghe et al., 2021), we select the XLM-RoBERTa (XLM-R) (Conneau et al., 2020) as the backbone of our model. XLM-R is a transformer-based masked language model pre-trained on large-scale multilingual corpus and demonstrates state-ofthe-art performance on multiple cross-lingual downstream tasks. As shown in Figure 5a, we concatenate the source sentence and the MT sentence together to make an input sample: $x_i =$ <s $>w_1^{\rm src},\ldots,w_m^{\rm src}</$ s><s $>w_1^{\rm mt},\ldots,w_n^{\rm mt}</$ s>, where m is the length of the source sentence (src) and n is the length of the MT sentence (mt). $\langle s \rangle$ and </s> are two special tokens to annotate the start and the end of the sentence in XLM-R, respectively.

For the j-th token $w_j^{\rm mt}$ in the MT sentence, we take the corresponding representation from XLM-R for binary classification to determine whether w_j belongs to good translation (OK) or contains translation error (BAD) and use the binary classification

³As a result, we do not include the sub-task of predicting gap tags in our experiments.

⁴Here, the BAD spans indicate the longest continuous tokens with BAD tags.

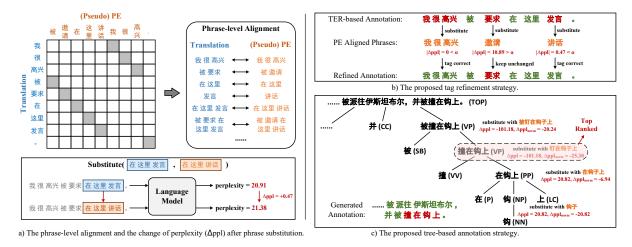


Figure 6: The proposed two tag correcting strategies: Tag Refinement strategy and Tree-based Annotation strategy.

loss to train the model:

$$s_{ij} = \sigma(\boldsymbol{w}^{\mathsf{T}} \mathsf{XLM-R}_{j}(\boldsymbol{x}_{i})) \tag{1}$$
$$\mathcal{L}_{ij} = -(y \cdot \log s_{ij} + (1 - y) \cdot \log(1 - s_{ij})) \tag{2}$$

where XLM-R_j(x_i) $\in \mathbb{R}^d$ (d is the hidden size of XLM-R) indicates the representation output by XLM-R corresponding to the token w_j^{mt} , σ is the sigmoid function, $\boldsymbol{w} \in \mathbb{R}^{d \times 1}$ is the linear layer for binary classification and y is the ground truth label.

3.2 Self-Supervised Pre-training Approach

The translation knowledge contained in the parallel corpus of MT is very helpful for the QE task. As a result, we can adopt the parallel corpus to build artificial tags for self-supervised pre-training on QE. As shown in Figure 5b, the parallel corpus is firstly split into the training and the test set. Then the NMT model is trained with the training split and is used to generate translations for all sentences in the test split. From this, a large number of triplets are obtained, each consisting of source, MT, and target sentences. Finally, the target sentence is regarded as the pseudo-PE from the MT sentence, and the TER toolkit is used to generate word-level OK | BAD tags based on the principle of minimal editing (shown in the bottom of Figure 1).

3.3 Tag Correcting Strategies

As we discussed before, the two issues of TER-based tags limit the performance improvement of the self-supervised pre-training when applied to the downstream translation error detection task. In this section, we introduce two tag correcting strategies, namely tag refinement and tree-based annotation, that target these issues and make the

TER-generated artificial QE tags more consistent with human judgment.

Tag Refinement Strategy. In response to the first issue (i.e., wrong annotations due to the synonym substitution or constituent reordering), we propose the tag refinement strategy, which corrects the false BAD tags to OK. Specifically, as shown in Figure 6a, we first generate the alignment between the MT sentence and the reference sentence (i.e., the pseudo-PE) using FastAlign⁵ (Dyer et al., 2013). Then we extract the phrase-to-phrase alignment through running the phrase extraction algorithm of NLTK⁶ (Bird, 2006). Once the phraselevel alignment is prepared, we substitute each BAD span with the corresponding aligned spans in the pseudo-PE and use the language model⁷ to calculate the change of the perplexity Δppl after this substitution.

If $|\Delta ppl| < \alpha$, where α is a hyper-parameter indicating the threshold, we regard that the substitution has little impact on the semantic and thus correct the BAD tags to OK. Otherwise, we regard the span does contain translation errors and keep the BAD tags unchanged (Figure 6b).

Tree-based Annotation Strategy. Human's direct annotation tends to annotate the *smallest* constituent that causes fatal translation errors *as a whole* (e.g., the whole words, phrases, clauses, etc.). However, TER-based annotations are often fragmented, with the whole mis-translations being split into multiple BAD spans because some stopwords are aligned and labeled as OK. Besides, the

⁵https://github.com/clab/fast_align
6https://github.com/nltk/nltk/blob/
develop/nltk/translate/phrase_based.py
7https://github.com/kpu/kenlm

Model	English-German (En-De)				English-Chinese (En-Zh)				
nivaei	MCC	F-OK	F-BAD	F-BAD-Span	MCC	F-OK	F-BAD	F-BAD-Span	
			Basel	ines					
FT on <i>HJQE</i> only	26.29	95.08	31.09	20.97	38.56	90.76	47.56	26.66	
PT (TER-based)	9.52	34.62	13.54	3.09	15.17	36.66	31.53	2.40	
+ FT on <i>HJQE</i>	24.82	94.65	29.82	18.52	39.09	91.29	47.04	25.93	
Pre-training only with tag correcting strategies (ours)									
PT w/ Tag Refinement	10.12*	49.33	14.32	3.62	19.36*	53.16	34.10	3.79	
PT w/ Tree-based Annotation	8.94	84.50	15.84	6.94	21.53*	59.21	35.54	6.32	
Pre-training with tag correcting strategies + fine-tuning on HJQE (ours)									
PT w/ Tag Refinement + FT	27.54*	94.21	35.25	21.13	40.35*	90.88	49.33	25.60	
PT w/ Tree-based Annotation + FT	27.67*	94.44	32.41	21.38	41.33*	91.22	49.82	27.21	

Table 2: The word-level QE performance on the test set of HJQE for two language pairs, En-De and En-Zh. PT indicates pre-training and FT indicates fine-tuning. Results are all reported by $\times 100$. The numbers with * indicate the significant improvement over the corresponding baseline with p < 0.05 under t-test (Semenick, 1990).

BAD spans are often not well-formed in linguistics (e.g., two adjacent words but are from two different phrases).

To address this issue, we propose the constituent tree-based annotation strategy. It can be regarded as an enhanced version of the tag refinement strategy that gets rid of the TER-based annotation. As shown in Figure 6c, we first generate the constituent tree for the MT sentences. Each internal node (i.e., the non-leaf node) in the constituent tree represents a well-formed phrase such as noun phrase (NP), verb phrase (VP), prepositional phrase (PP), etc. For each node, we substitute it with the corresponding aligned phrase in the pseudo-PE. Then we still use the change of the perplexity Δppl to indicate whether the substitution of this phrase improves the fluency of the whole translation.

To only annotate the smallest constituents that exactly contain translation errors, we normalize Δppl by the number of words in the phrase and use this value to sort all internal nodes in the constituent tree: $\Delta ppl_{\rm norm} = \frac{\Delta ppl}{r-l+1},$ where l and r indicates the left and right position of the phrase, respectively. The words of a constituent node are integrally labeled as BAD only if $|\Delta ppl_{\rm norm}|<\beta$ as well as there is no overlap with nodes that are higher ranked. β is a hyperparameter indicating the threshold.

4 Experiments

Datasets. To verify the effectiveness of our proposed self-supervised pre-training approach with tag correcting strategies on detecting translation errors, we conduct experiments on both

HJQE and MLQE-PE (Fomicheva et al., 2020) datasets. MLQE-PE is the official dataset used in the WMT20 QE shared task (Specia et al., 2020), and HJQE is our collected dataset with word-level annotations of translation errors. Note that MLQE-PE and HJQE share the same source and MT sentences, thus they have exactly the same number of samples. We show the detailed statistics in Table 1. For the pre-training, we use the parallel dataset provided in the WMT20 QE shared task to generate the artificial OE dataset.

Baselines. To confirm the effectiveness of our proposed self-supervised pre-training approach with tag correcting strategies, we mainly select two baselines for comparison. In the one, we do not use the pre-training, but only fine-tune XLM-R on the training set of *HJQE*. In the other, we pre-train the model on the TER-based artificial QE dataset and then fine-tune it on the training set of *HJQE*.

Evaluation. Following WMT20 QE shared task (Specia et al., 2020), we use Matthews Correlation Coefficient (MCC) as the main metric and also provide the F1 score (F) for OK, BAD and BAD spans.⁹

4.1 Main Results

The results are shown in Table 2. We can observe that the TER-based pre-training only brings very limited performance gain or even degrade the performance when compared to the "FT on *HJQE* only" setting (-1.47 for En-De and +0.53 for En-Zh). It suggests that the inconsistency between TER-based and human's annotations leads to the limited effect of pre-training. However, when applying the tag correcting strategies to the pre-

⁸https://stanfordnlp.github.io/stanza

⁹Please refer to Appendix A for implementation details.

training dataset, the improvement is much more significant (+2.85 for En-De and +2.24 for En-Zh), indicating that the tag correcting strategies mitigate such inconsistency, improving the effect of pre-training. On the other hand, when only the pre-training is applied, the tag correcting strategies can also improve the performance. It shows our approach can also be applied to the unsupervised setting, where no human-annotated dataset is available for fine-tuning.

Tag Refinement v.s. Tree-based Annotation. When comparing two tag correcting strategies, we find the tree-based annotation strategy is generally superior to the tag refinement strategy, especially for En-Zh. The MCC improves from 19.36 to 21.53 under the *pre-training only* setting and improves from 40.35 to 41.33 under the *pre-training then fine-tuning* setting. This is probably because the tag refinement strategy still requires the TER-based annotation and fixes based on it, while the tree-based annotation strategy actively selects the well-formed constituents to apply phrase substitution and gets rid of the TER-based annotation.

Span-level Metric. Through the span-level metric (F-BAD-Span), we want to measure the unity and consistency of the model's prediction against human judgment. From Table 2, we find our models with tag correcting strategies also show higher F1 score on BAD spans (from 26.66 to 27.21 for En-Zh), while TER-based pre-training even do harm to this metric (from 26.66 to 25.93 for En-Zh). This phenomenon also confirms the aforementioned fragmented issue of TER-based annotations, and our tag correcting strategies, instead, improve the span-level metric by alleviating this issue.

4.2 Analysis

Comparison to results on MLQE-PE. To demonstrate the difference between the MLQE-PE (TERgenerated tags) and our *HJQE* datasets, and analyze how the pre-training and fine-tuning influence the results on both datasets, we compare the performance of different models on MLQE-PE and *HJQE* respectively. The results for En-Zh are shown in Table 3.

When comparing results in each group, we find that fine-tuning on the training set identical to the evaluation set is necessary for achieving high performance. Otherwise, fine-tuning provides marginal improvement (e.g., fine-tuning on MLQE-PE and evaluating on *HJQE*) or even degrades the

Evaluate on \rightarrow	N	/ILQE-P	HJQE					
Fine-tune on ↓	MCC*	MCC	F-BAD	MCC	F-BAD			
WMT20's best	59.28	-	-	-	-			
No pre-training (fine-tuning only)								
MLQE-PE	58.21	46.81	75.02	22.49	34.34			
HJQE	49.77	23.68	36.10	45.76	53.77			
TER-based pre-training								
w/o fine-tune	56.51	33.58	73.85	11.38	27.41			
MLQE-PE	61.85	53.25	78.69	21.93	33.75			
HJQE	41.39	29.19	42.97	47.34	55.43			
Pre-training with tag refinement								
w/o fine-tune	55.03	28.89	70.73	18.83	31.39			
MLQE-PE	61.35	48.24	77.17	21.85	33.31			
HJQE	39.56	25.06	67.40	47.61	55.22			
Pre-training with tree-based annotation								
w/o fine-tune	55.21	26.79	68.11	20.98	32.84			
MLQE-PE	60.92	48.58	76.18	22.34	34.13			
HJQE	40.30	26.22	39.50	48.14	56.02			

Table 3: Performance comparison for En-Zh with different fine-tuning and evaluation settings. Since the test labels of MLQE-PE are not publicly available, we report the results on the validation set of both datasets. MCC* indicates the MCC score considering both the target tokens and the target gaps.

performance (e.g., fine-tuning on *HJQE* and evaluating on MLQE-PE). This reveals the difference in data distribution between *HJQE* and MLQE-PE. Besides, we note that our best model on MLQE-PE outperforms WMT20's best model (61.85 v.s. 59.28) using the same MCC* metric, showing the strength of our model, even under the TER-based setting.

On the other hand, we compare the performance gain of different pre-training strategies. When evaluating on MLQE-PE, the TER-based pre-training brings higher performance gain (+6.44) than pre-training with two proposed tag correcting strategies (+1.43 and +1.77). While when evaluating on *HJQE*, the case is opposite, with the TER-based pre-training bringing lower performance gain (+1.58) than tag refinement (+1.85) and tree-based annotation (+2.38) strategies. In conclusion, the pre-training always brings performance gain, no matter evaluated on MLQE-PE or *HJQE*. However, the optimal strategy depends on the consistency between the pre-training dataset and the downstream evaluation task.

Human Evaluation. To evaluate and compare the models pre-trained on TER-based tags and corrected tags more objectively, human evaluation is conducted for both models. For En-Zh and En-De, we randomly select 100 samples (the source and MT sentences) from the validation set and use two models to predict word-level OK or BAD tags for

Scores	En	-De	En-Zh			
Scores	TER	Ours	TER	Ours		
1 (terrible)	3	1	5	0		
2 (bad)	36	16	34	6		
3 (neutral)	34	20	29	21		
4 (good)	26	61	24	59		
5 (excellent)	1	2	8	14		
Average score:	2.86	3.47	2.96 3.81			
% Ours \geq TER:	89	9%	91%			

Table 4: The results of human evaluation. We select the best-performed model fine-tuned on MLQE-PE and *HJQE* respectively.

them. Then, we ask human translators to give a score for each prediction, between 1 and 5, where 1 indicates the predicted tags are fully wrong, and 5 indicates the tags are fully correct.

Table 4 shows the results. We can see that the model pre-trained on corrected tags (Ours) achieves higher human evaluation scores than that pre-trained on TER-based tags on average. For about 90% of samples, the prediction of the model pre-trained on corrected dataset can outperform or tie with the prediction of the model pre-trained on TER-based dataset. The results of human evaluation show that *HJQE* is more consistent with human judgement.

5 Related Work

Early approaches on QE, such as QuEst (Specia et al., 2013) and QuEst++ (Specia et al., 2015), mainly pay attention to the feature engineering. They aggregate various features and feed them to the machine learning algorithms for classification or regression. Kim et al. (2017) first propose the neural-based QE approach, called Predictor-Estimator. They first pre-train an RNN-based predictor on the large-scale parallel corpus that predicts the target word given its context and the source sentence. Then, they extract the features from the pre-trained predictor and use them to train the estimator for the QE task. This model achieves the best performance on the WMT17 QE shard task. After that, many variants of Predictor-Estimator are proposed (Fan et al., 2019; Moura et al., 2020; Cui et al., 2021). Among them, Bilingual Expert (Fan et al., 2019) replaces RNN with multi-layer transformers as the architecture of the predictor, and proposes the 4-dimension mismatching feature for each token. It achieves the best performance on WMT18 QE shared task. The Unbabel team also

releases an open-source framework for QE, called OpenKiwi (Kepler et al., 2019), that implements the most popular QE models with configurable architecture.

Recently, with the development of pre-trained language models, many works select the crosslingual language model XLM-RoBERTa (Conneau et al., 2020) as the backbone (Ranasinghe et al., 2020; Lee, 2020; Moura et al., 2020; Rubino and Sumita, 2020; Ranasinghe et al., 2021; Zhao et al., 2021). Many works also explore the joint learning or transfer learning of the multilingual QE task (i.e., on many language pairs) (Sun et al., 2020; Ranasinghe et al., 2020, 2021). Meanwhile, on the word-level OE, Fomicheva et al. (2021) propose a shared task with the new-collected dataset on explainable QE, aiming to provide word-level hints for sentence-level QE score. Freitag et al. (2021) also study multidimensional human evaluation for MT and collect a large-scale dataset.

The QE model can be applied to the Computer-Assisted Translation (CAT) system together with other models like translation suggestion (TS) or automatic post-edit (APE). Wang et al. (2020) and Lee et al. (2021) use the QE model to identify which parts of the machine translations need to be correct, and the TS (Yang et al., 2021) also needs the QE model to determine error spans before giving translation suggestions.

6 Conclusion

In this paper, we focus on the task of word-level QE in machine translation and target the inconsistency issues between the TER-based QE dataset and human judgment. We first collect and release a benchmark dataset called HJQE that reflects the human judgement on the translation errors in MT sentences. Besides, we propose the self-supervised pre-training approach with two tag correcting strategies, which makes the TER-based annotations closer to the human judgement and improves the final performance on the proposed benchmark dataset *HJQE*. We conduct thorough experiments and analyses, demonstrating the necessity of our proposed dataset and the effectiveness of our proposed approach. Our future directions include improving the performance of phraselevel alignment, introducing phrase-level semantic matching, and applying data augmentation. We hope our work will provide a new perspective for future researches on quality estimation.

Broader Impacts

Quality estimation often serves as a postprocessing module in recent commercial machine translation systems. It can be used to indicate the overall translation quality or detect the specific translation errors in the sentences. This work focuses on the direct annotation of translation errors, training the model to fit the human judgment at the word level. To do this, we collect a new QE dataset and propose tag correcting strategies to force the TER-based artificial dataset used in the pre-training phase closer to human judgment. When applying our approach, the users should pay special attention to the following: a) The data source of HJQE is Wikipedia, so our model should perform well on a similar domain but may perform poorly on other irrelevant domains. b) Since our approach is still data-driven, the data (as well as the pre-training parallel dataset) should be ethical and unbiased, or unexpected problems may arise. c) The proposed tag correcting strategies work well on En-De and En-Zh, but do not necessarily applicable to other language pairs since the characteristics among target languages are different. d) Since the system is neural-based, the interpretability is limited. It can still mistakenly annotate some forbidden or sensitive words to OK and cause unexpected accidents.

References

- Steven Bird. 2006. Nltk: the natural language toolkit. In *Proceedings of the COLING/ACL 2006 Interactive Presentation Sessions*, pages 69–72.
- Wanxiang Che, Zhenghua Li, and Ting Liu. 2010. Ltp: A chinese language technology platform. In *Coling* 2010: Demonstrations, pages 13–16.
- Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Édouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. 2020. Unsupervised cross-lingual representation learning at scale. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 8440–8451.
- Qu Cui, Shujian Huang, Jiahuan Li, Xiang Geng, Zaixiang Zheng, Guoping Huang, and Jiajun Chen. 2021. Directqe: Direct pretraining for machine translation quality estimation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 12719–12727.
- Chris Dyer, Victor Chahuneau, and Noah A Smith. 2013. A simple, fast, and effective reparameterization of ibm model 2. In *Proceedings of the 2013*

- Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, pages 644–648.
- Kai Fan, Jiayi Wang, Bo Li, Fengming Zhou, Boxing Chen, and Luo Si. 2019. "bilingual expert" can find translation errors. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 6367–6374.
- Marina Fomicheva, Piyawat Lertvittayakumjorn, Wei Zhao, Steffen Eger, and Yang Gao. 2021. The eval4nlp shared task on explainable quality estimation: Overview and results. In *Proceedings of the 2nd Workshop on Evaluation and Comparison of NLP Systems*, pages 165–178.
- Marina Fomicheva, Shuo Sun, Erick Fonseca, Frédéric Blain, Vishrav Chaudhary, Francisco Guzmán, Nina Lopatina, Lucia Specia, and André FT Martins. 2020. Mlqe-pe: A multilingual quality estimation and post-editing dataset. *arXiv preprint arXiv:2010.04480*.
- Markus Freitag, George Foster, David Grangier, Viresh Ratnakar, Qijun Tan, and Wolfgang Macherey. 2021. Experts, errors, and context: A large-scale study of human evaluation for machine translation. *arXiv* preprint arXiv:2104.14478.
- Kenneth Heafield. 2011. Kenlm: Faster and smaller language model queries. In *Proceedings of the sixth workshop on statistical machine translation*, pages 187–197.
- Fabio Kepler, Jonay Trénous, Marcos Treviso, Miguel Vera, and André FT Martins. 2019. Openkiwi: An open source framework for quality estimation. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pages 117–122.
- Hyun Kim, Hun-Young Jung, Hongseok Kwon, Jong-Hyeok Lee, and Seung-Hoon Na. 2017. Predictorestimator: Neural quality estimation based on target word prediction for machine translation. *ACM Transactions on Asian and Low-Resource Language Information Processing (TALLIP)*, 17(1):1–22.
- Diederik P Kingma and Jimmy Ba. 2014. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.
- Dongjun Lee. 2020. Two-phase cross-lingual language model fine-tuning for machine translation quality estimation. In *Proceedings of the Fifth Conference on Machine Translation*, pages 1024–1028.
- Dongjun Lee, Junhyeong Ahn, Heesoo Park, and Jaemin Jo. 2021. Intellicat: Intelligent machine translation post-editing with quality estimation and translation suggestion. *arXiv preprint arXiv:2105.12172*.

- Joao Moura, Miguel Vera, Daan van Stigt, Fabio Kepler, and André FT Martins. 2020. Ist-unbabel participation in the wmt20 quality estimation shared task. In *Proceedings of the Fifth Conference on Machine Translation*, pages 1029–1036.
- Peng Qi, Yuhao Zhang, Yuhui Zhang, Jason Bolton, and Christopher D Manning. 2020. Stanza: A python natural language processing toolkit for many human languages. arXiv preprint arXiv:2003.07082.
- Tharindu Ranasinghe, Constantin Orasan, and Ruslan Mitkov. 2020. Transquest: Translation quality estimation with cross-lingual transformers. In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 5070–5081.
- Tharindu Ranasinghe, Constantin Orasan, and Ruslan Mitkov. 2021. An exploratory analysis of multilingual word-level quality estimation with cross-lingual transformers. *arXiv preprint arXiv:2106.00143*.
- Raphael Rubino and Eiichiro Sumita. 2020. Intermediate self-supervised learning for machine translation quality estimation. In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 4355–4360.
- Doug Semenick. 1990. Tests and measurements: The t-test. *Strength & Conditioning Journal*, 12(1):36–37.
- Matthew Snover, Bonnie Dorr, Richard Schwartz, Linnea Micciulla, and John Makhoul. 2006. A study of translation edit rate with targeted human annotation. In *Proceedings of the 7th Conference of the Association for Machine Translation in the Americas: Technical Papers*, pages 223–231.
- Lucia Specia, Frédéric Blain, Marina Fomicheva, Erick Fonseca, Vishrav Chaudhary, Francisco Guzmán, and André F. T. Martins. 2020. Findings of the WMT 2020 shared task on quality estimation. In *Proceedings of the Fifth Conference on Machine Translation*, pages 743–764, Online. Association for Computational Linguistics.
- Lucia Specia, Gustavo Paetzold, and Carolina Scarton. 2015. Multi-level translation quality prediction with quest++. In *Proceedings of ACL-IJCNLP 2015 System Demonstrations*, pages 115–120.
- Lucia Specia, Kashif Shah, José GC De Souza, and Trevor Cohn. 2013. Quest-a translation quality estimation framework. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pages 79–84.
- Shuo Sun, Marina Fomicheva, Frédéric Blain, Vishrav Chaudhary, Ahmed El-Kishky, Adithya Renduchintala, Francisco Guzmán, and Lucia Specia. 2020. An exploratory study on multilingual quality estimation. In *Proceedings of the 1st Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 10th International Joint*

- Conference on Natural Language Processing, pages 366–377.
- Yi-Lin Tuan, Ahmed El-Kishky, Adithya Renduchintala, Vishrav Chaudhary, Francisco Guzmán, and Lucia Specia. 2021. Quality estimation without humanlabeled data. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 619–625, Online. Association for Computational Linguistics.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. In *Advances in neural information processing systems*, pages 5998–6008.
- Ke Wang, Jiayi Wang, Niyu Ge, Yangbin Shi, Yu Zhao, and Kai Fan. 2020. Computer assisted translation with neural quality estimation and auotmatic postediting. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: Findings*, pages 2175–2186.
- Shufan Wang, Laure Thompson, and Mohit Iyyer. 2021. Phrase-bert: Improved phrase embeddings from bert with an application to corpus exploration. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 10837–10851.
- Zhen Yang, Yingxue Zhang, Ernan Li, Fandong Meng, and Jie Zhou. 2021. Wets: A benchmark for translation suggestion. *arXiv preprint arXiv:2110.05151*.
- Mingjun Zhao, Haijiang Wu, Di Niu, Zixuan Wang, and Xiaoli Wang. 2021. Verdi: Quality estimation and error detection for bilingual corpora. In *Pro*ceedings of the Web Conference 2021, pages 3023– 3031.

A Implementation Details

Our implementation of QE model is based on an open-source framework, OpenKiwi¹⁰ (Kepler et al., 2019). We use the large-sized XLM-R model and obtain it from hugging-face's library¹¹. We use the KenLM¹² (Heafield, 2011) to train the language model on all target sentences in the parallel corpus and calculate the perplexity of the given sentence. For the tree-based annotation strategy, we obtain the constituent tree through LTP¹³ (Che et al., 2010) for Chinese and through Stanza¹⁴ (Qi et al., 2020) for German. We set α to 1.0 and β to -3.0 in our tag correcting strategies based on the case studies and empirical judgment. In the preprocessing phase, we filter out parallel samples that are too long or too short, and only reserve sentences with 10-100 tokens.

We pre-train the model on 8 NVIDIA Tesla V100 (32GB) GPUs for two epochs, with the batch size set to 8 for each GPU. Then we fine-tune the model on a single NVIDIA Tesla V100 (32GB) GPU for up to 10 epochs, with the batch size set to 8 as well. Early stopping is used in the fine-tuning phase, with the patience set to 20. We evaluate the model every 10% steps in one epoch. The pre-training often takes more than 15 hours and the fine-tuning takes 1 or 2 hours. We use Adam (Kingma and Ba, 2014) to optimize the model with the learning rate set to 5e-6 in both the pre-training and fine-tuning phases. For all hyperparameters in our experiments, we manually tune them on the validation set of *HJQE*.

B Main Results on the Validation Set

In Table 5, we also report the main results on the validation set of *HJQE*.

C Case Study

In Figure 7, we show some cases from the validation set of English-Chinese language pair. From the examples, we can see that the TER-based model (noted as PE Effort Prediction) often annotates wrong BAD spans and is far from human judgment. For the first example, the MT sentence correctly

reflects the meaning of the source sentence, and the PE is just a paraphrase of the MT sentence. Our model correctly annotates all words as OK, while TER-based one still annotates many BAD words. For the second example, the key issue is the translation of "unifies" in Chinese. Though "统一" is the direct translation of "unifies" in Chinese, it can not express the meaning of winning two titles in Chinese context. And our model precisely annotated the "统一了" in the MT sentence as BAD. For the third example, the MT model fails to translate the "parsley" and the "sumac" to "欧芹" and "盐 肤木" in Chinese, since they are very rare words. While the TER-based model mistakenly predicts long BAD spans, our model precisely identities both mistranslated parts in the MT sentence.

D Limitation and Discussion

We analyze some samples that are corrected by our tag correcting strategies and find a few bad cases. These are mainly because of the following:

1) There is noise from the parallel corpus (i.e., the source sentence and the target sentence are not well aligned).

2) The alignment generated by FastAlign contains unexpected errors, making some entries in the phrase-level alignments are missing or misaligned.

3) The scores given by KenLM (through the change of the perplexity after the phrase substitution) are sometimes not consistent with human judgment.

We also propose some possible solutions in response to the above problems as our future exploration direction. For the noise in the parallel corpus, we can use parallel corpus filtering methods that filter out samples with low confidence. We can also apply the data augmentation methods that expand the corpus based on the clean parallel corpus. For the errors by FastAlign, we may use a more accurate alignment model. For the scoring, we may introduce the neural-based phrase-level semantic matching model (e.g., Phrase-BERT (Wang et al., 2021)) instead of the KenLM.

Model	English-German (En-De)				English-Chinese (En-Zh)				
Model	MCC	C F-OK F-BAD F-BAD-Span		MCC	F-OK	F-BAD	F-BAD-Span		
Baselines									
FT on <i>HJQE</i> only	34.69	94.28	40.38	28.65	45.76	91.96	53.77	29.84	
PT (TER-based)	13.13	37.30	18.80	4.72	11.38	25.91	27.41	2.16	
+ FT on <i>HJQE</i>	35.02	94.00	40.86	26.68	47.34	91.30	55.43	28.53	
With tag correcting strategies (ours)									
PT w/ Tag Refinement	13.26	52.43	19.78	6.42	18.83	53.29	31.39	3.48	
+ FT on <i>HJQE</i>	37.70	94.08	43.32	30.83	47.61	92.39	55.22	28.33	
PT w/ Tree-based Annotation	13.92	84.79	22.75	9.64	20.98	59.32	32.84	6.53	
+ FT on <i>HJQE</i>	37.03	94.46	42.54	31.21	48.14	91.88	56.02	28.17	
PT w/ Both	13.12	39.68	18.94	5.26	21.39	56.76	32.74	5.72	
+ FT on <i>HJQE</i>	38.90	94.44	44.35	32.21	48.71	90.74	56.47	25.51	

Table 5: The word-level QE performance on the validation set of *HJQE* for two language pairs, En-De and En-Zh. PT indicates pre-training and FT indicates fine-tuning.

```
Source: To win, a wrestler must strip their opponent's tuxedo off.
```

MT: 要 想 获胜, 摔跤 运动员 必须 把 对手 的 礼服 脱下来.

MT Back: To win, the wrestler had to take his opponent's dress off.

PE: 要 赢得 胜利 , 摔跤 运动员 必须 脱掉 对手 的 燕尾服 。

PE Back: To win the victory, the wrestler had to remove his opponent's tuxedo.

TER-based: 要想 获胜, 摔跤 运动员 必须 把 对手 的 礼服 脱下来.

Ours: 要 想 获胜, 摔跤 运动员 必须 把 对手 的 礼服 脱下来.

Source: April 28 Juan Díaz unifies the WBA and WBO Lightweight titles after defeating Acelino Freitas.

MT: 4 月 28 日 , 胡安 · 迪亚斯 在 击败 阿 切利 诺 · 弗雷 塔斯 后 统一 了 WBA 和 WBO 轻量级 冠军 .

MT Back: On April 28, Juan Díaz Unified the WBA and WBO lightweight titles after defeating Acelino Freitas.

PE: 4 月 28 日 , Juan Díaz 在 击败 Acelino Freitas 之后 , 将 W 世界 拳击 协会 和 世界 拳击 组织 的 轻量级 冠军 揽于 一身 。

PE Back: On April 28, Juan Díaz won both the WBA and WBO lightweight titles after defeating Acelino Freitas.

TER-based: 4 月 28 日 , 胡安 · 迪亚斯 在 击败 阿 切利 诺 · 弗雷 塔斯 后 统一 了 WBA 和 WBO 轻量级 冠军 .

Ours: 4月 28日, 胡安·迪亚斯 在 击败 阿 切利 诺·弗雷 塔斯 后 统一 了 WBA 和 WBO 轻量级 冠军.

Source: Fattoush is a combination of toasted bread pieces and parsley with chopped cucumbers, radishes, tomatoes and flavored by sumac.

MT: 法杜什是 烤面包片 和 帕斯 莱 与 切碎 的 黄瓜 、 萝卜 、 西红柿 、 和 洋葱 以及 香味 的 消耗品 的 组合 。

MT Back: Fadush is a combination of toast and pasai with chopped cucumbers, radishes, tomatoes and onions and scented consumables.

PE: Fattoush 是 烤面包片 和 欧芹 与 切碎 的 黄瓜 , 萝卜 , 西红柿 和 葱 的 组合 , 并 以 盐肤木 调味 。

PE Back: Fattoush is a combination of toast and parsley with chopped cucumbers, radishes, tomatoes and scallions, seasoned with rhus salt.

TER-based: 法杜什是 烤面包片 和 帕斯 莱 与 切碎 的 黄瓜 imes 萝卜 imes 西红柿 imes 和 洋葱 以及 香味 的 消耗品 的 组合 。

Ours: 法杜什是 烤面包片 和 帕斯 莱 与 切碎 的 黄瓜 、 萝卜 、 西红柿 、 和 洋葱 以及 香味 的 消耗品 的 组合 。

Figure 7: Examples of word-level QE from the validation set of English-Chinese language pair.